

# Deqing Fu

Email: deqingfu@usc.edu  
Website: www.fudeq.ing

Address: Ginsburg Hall, USC, Los Angeles, CA 90089  
Office: GCS 403E

Education	<b>University of Southern California</b> <i>Ph.D. in Computer Science</i> Advisors: Vatsal Sharan & Robin Jia	2022 - Present
	<b>University of Chicago</b> <i>M.S. in Statistics</i>	2020 - 2022
	<b>University of Chicago</b> <i>B.S. in Mathematics with Honors</i> <i>B.S. in Computer Science with Honors</i> <i>B.A. in Statistics</i>	2016 - 2020
Experience	<b>Simons Institute at UC Berkeley</b> <i>Visiting Graduate Student</i> Program: Modern Paradigms in Generalization	Fall 2024
	<b>Research Scientist Intern</b> $\rightarrow$ <b>Student Researcher</b> , Meta Team: Llama Team @ Meta GenAI. Host: Lawrence Chen; Peers: Tong Xiao, Rui Wang, and Pengchuan Zhang.	May – Dec., 2024
	<b>Software Engineering Intern</b> , Google Teams: Image Search (Lens); Knowledge Engine; Bard Multilingual.	Summer 2021, 2022, 2023
	<b>Summer Student Researcher</b> , Argonne National Laboratory Hosts: Paul Hovland & Sri Hari Krishna Narayana.	Summer 2017
Publications  * denotes equal contribution	<b>TLDR: Token-Level Detective Reward Model for Large Vision Language Models</b> Deqing Fu, Tong Xiao, Rui Wang, Wang Zhu, Pengchuan Zhang, Guan Pang, Robin Jia, Lawrence Chen	ICLR 2025
	<b>Transformers Learn Low Sensitivity Functions: Investigations and Implications</b> Bhavya Vasudeva*, Deqing Fu*, Tianyi Zhou, Elliot Kau, You-Qi Huang, Vatsal Sharan	ICLR 2025
	<b>DeLLMa: Decision Making Under Uncertainty with Large Language Models</b> Ollie Liu*, Deqing Fu*, Dani Yogatama, Willie Neiswanger <i>ICLR 2025 Spotlight (Top 5.1%)</i>	ICLR 2025
	<b>DreamSync: Aligning Text-to-Image Generation with Image Understanding Feedback</b> Jiao Sun*, Deqing Fu*, Yushi Hu*, Su Wang, Royi Rassin, Da-Cheng Juan, Dana Alon, Charles Herrmann, Sjoerd van Steenkiste, Ranjay Krishna, Cyrus Rashtchian	NAACL 2025
	<b>Transformers Learn to Achieve Second-Order Convergence Rates for In-Context Linear Regression</b> Deqing Fu, Tian-Qi Chen, Robin Jia, Vatsal Sharan <i>SoCalNLP Symposium 2023 Best Paper Award</i>	NeurIPS 2024

	<b>Pre-trained Large Language Models Use Fourier Features to Compute Addition</b> Tianyi Zhou, <a href="#">Deqing Fu</a> , Vatsal Sharan, Robin Jia	NeurIPS 2024
	<b>IsoBench: Benchmarking Foundation Models on Isomorphic Representations</b> <a href="#">Deqing Fu</a> <sup>*</sup> , Ruohao Guo <sup>*</sup> , Ghazal Khalighinejad <sup>*</sup> , Ollie Liu <sup>*</sup> , Bhuwan Dhingra, Dani Yogatama, Robin Jia, Willie Neiswanger <a href="#">Huggingface</a> 🗳️: [ Dataset ]	COLM 2024
	<b>SCENE: Self-Labeled Counterfactuals for Extrapolating to Negative Examples</b> <a href="#">Deqing Fu</a> , Ameya Godbole, Robin Jia	EMNLP 2023
	<b>Harnessing the Conditioning Sensorium for Improved Image Translation</b> Cooper Nederhood, Nicholas Kolkin, <a href="#">Deqing Fu</a> , Jason Salavon	ICCV 2021
Preprints	<b>Zebra-CoT: A Dataset for Interleaved Vision Language Reasoning</b> Ang Li <sup>*</sup> , Charles L. Wang <sup>*</sup> , Kaiyu Yue <sup>*</sup> , Zikui Cai <sup>*</sup> , Ollie Liu <sup>*</sup> , <a href="#">Deqing Fu</a> <sup>*</sup> , Peng Guo <sup>*</sup> , Wang Bill Zhu <sup>*</sup> , Vatsal Sharan, Robin Jia, Willie Neiswanger, Furong Huang, Tom Goldstein, Micah Goldblum <a href="#">Huggingface</a> 🗳️: [ Dataset ]	Under Review
	<b>Textual Steering Vectors Can Improve Visual Understanding in Multimodal Large Language Models</b> Woody Haosheng Gan <sup>*</sup> , <a href="#">Deqing Fu</a> <sup>*</sup> , Julian Asilis <sup>*</sup> , Ollie Liu <sup>*</sup> , Dani Yogatama, Vatsal Sharan, Robin Jia, Willie Neiswanger	Under Review
	<b>FoNE: Precise Single-Token Number Embeddings via Fourier Features</b> Tianyi Zhou, <a href="#">Deqing Fu</a> , Mahdi Soltanolkotabi, Robin Jia, Vatsal Sharan	Under Review
	<b>VisualLens: Personalization through Visual History</b> Wang Bill Zhu, <a href="#">Deqing Fu</a> , Kai Sun, Yi Lu, Zhaojiang Lin, Seungwhan Moon, Kanika Narang, Mustafa Canim, Yue Liu, Anuj Kumar, Xin Luna Dong	Under Review
Awards	<b>Viterbi Best Research Assistant Award</b> ( <i>one per department</i> ) <b>SoCalNLP Symposium 2023 Best Paper Award</b> <b>Provost's Fellowship</b> <b>Susanne H. Rudolph Scholarship</b> <b>Liew Family College Research Fellows Fund</b> <b>Dean's List</b>	USC, 2025 SoCalNLP, 2023 USC, 2022 UChicago, 2020 UChicago, 2020 UChicago, 2016-2020
Invited Talks	<i>Closing the Modality Gap: Benchmarking and Improving Visual Understanding in Multimodal LLMs</i> Stanford NLP Seminar  <i>Algorithmic Perspectives on Understanding Transformers</i> UChicago/TTIC NLP Seminar	May 22, 2025  May 08, 2025

	<i>Tutorial on the interplay between Theory and Interpretability</i> USC ML/NLP Joint Seminar	Apr 25, 2025
	<i>Algorithmic Perspectives on Understanding Transformers</i> Duke NLP Seminar	Feb 12, 2025
	<i>Guest Lecture on Interpretability and AI Safety</i> DSO 575: Driving Business Transformation with GenAI	Oct 24, 2024
	<i>Simplicity Bias of Transformers to Learn Low Sensitivity Functions</i> USC NLP Lunch	Feb 22, 2024
	<i>Aligning Text-to-Image Generation with Image Understanding Feedback</i> USC NLP Lunch	Oct 19, 2023
	<i>Transformers Learn to Achieve Second-Order Convergence for In-Context Linear Regression</i> USC NLP Lunch	Mar 29, 2023
<b>Services</b>	<b>Reviewer</b> for ICLR 2025, NeurIPS (2024, 2025), ICML 2025, COLM 2024, EMNLP (2023, 2024), ACL ARR 2024, ECCV 2022, CVPR (2023, 2024, 2025)	
	<b>Organizer</b> of USC ML/NLP Joint Reading Group	Jan 2025 – Present
	<b>Organizer</b> of USC AI Safety Club	Jan 2024 – Jan 2025
<b>Students Supervised</b>	<b>Master's and Undergraduate Students</b>	
	Gustavo Carvalho (USC MS)	Feb 2024 – Present
	Qilin Ye (USC UG $\rightarrow$ Duke MS)	Sep 2023 – Present
	Woody Haosheng Gan (USC UG)	Sep 2023 – Present
	You-Qi Huang (USC UG)	Sep 2023 – Mar 2024
	Elliot Kau (USC MS)	Sep 2023 – Mar 2024
	Tian-Qi Chen (USC MS)	Mar 2023 – Sep 2023
<b>Teaching Experience</b>	<b>Teaching Assistant</b> , CSCI 467: Introduction to Machine Learning University of Southern California, Los Angeles, CA	
		Spring 2025
	<b>Teaching Assistant</b> , CSCI 566: Deep Learning and its Applications University of Southern California, Los Angeles, CA	
		Spring 2023
	<b>Teaching Assistant</b> , BUSN 41301: Statistical Insight Booth School of Business, University of Chicago, Chicago, IL	
		Fall 2021
	<b>Course Assistant</b> , CMSC 31230: Fundamentals of Deep Learning University of Chicago, Chicago, IL	
		Fall 2021
	<b>Course Assistant</b> , MATH 20300: Analysis in $\mathbb{R}^n$ - I University of Chicago, Chicago, IL	
		Spring 2018
<b>Skills</b>	C/C++, Python, $\LaTeX$ , MATLAB, R, Java, CUDA PyTorch, TensorFlow, Keras, JAX, HuggingFace	